

Working with Databases

Relational Databases (case MySQL)

SQL Language

JDBC

This week

- You will learn to understand basic concepts of relational databases
- You will learn SQL-language (to sufficient level)
- You will learn how to use JDBC and Spring Data JDBC to connect to a database from Java Application
- And we will also take a brief look to NoSQL-databases (MongoDB) and other data formats commonly used

Relational Databases

Overview

RDBMS

- On the background there is the server machine that has the database product installed
 - The product may be Oracle, IBM DB2, Microsoft SQL Server, MySQL, PostgreSQL....
 - Products are often called DBMS or RDBMS (Database Management System, Relational Database Management System)
 - Sometimes the term Database Engine is also used to refer to DBMS
- Later on, we start developing server applications. Then the DBMS may be installed on
 - The same machine the application runs (rather rare case)
 - Behind the network on separate server
 - Big installations on server cluster
 - Several servers holding exactly the same applications and configurations to offer better response times and provide fault tolerance
- Developers may have the DBMS installed on their development machine
 - Or they may operate on development server over the network

Table

- Technically the table is the basic unit of storage under the DBMS
 - Think of Excel spreadsheet having columns and rows, each column also having a title
- Logically the table is also the basic relation
 - Individual items on each row have a relation to each other, together they describe some kind of data entity
 - On person table each row holds data items that belong to single person
- Most RDBMS products group tables to logical databases
 - Single application often operates to a single logical database

SQL Language

- SQL Language is used to operate on RDBMS
 - Standard even though each RDBMS has its own dialect
- Sometimes a distinction is made between
 - DDL – Data Definition language
 - Statements used to create and modify tables
 - DML – Data Manipulation Language
 - Statements used to maintain and query data in the database
- DML gives you CRUD
 - Create
 - Review
 - Update
 - Delete

Transaction

- SQL statements are always run by the database engine within a transaction
 - It may contain single sql statement
 - Or group of sql statements
- During the transaction the engine ensures that no other user accesses the same data
 - Someone writes while you are reading => you get partially written data, this never happens with databases
 - Database engine “locks” the data you access
 - No one can write the data while you read, no one read while you write
 - Sometimes if complex operations are executed this may result in dead-locks, no one can read or write....
- Transactions are ACID
 - Atomicity
 - Consistency
 - Isolation
 - Durability

A few more terms

- OLTP – Online Transaction Processing
 - The typical database –based system where a user manipulates data
 - Enter data, queries some listings etc
 - This is what we will be implementing
- OLAP – Online Analytical Processing
 - Database constructs are designed especially for reporting
 - Often “multidimensional” you can query sales report
 - By sales person
 - By sales area
 - By fiscal quarter
 - By customer group
 - Or any combination of above
- DW – Datawarehouse
 - A huge storage of business data
 - Often a source of data for OLAP system

Designing a database

- You should start with a conceptual model
 - What data entities are manipulated: Person, Company....
 - These will become tables of database
 - Often you will have similar classes in your application and the fields of those classes translate to columns in the tables
- For a good database design there are quite a few guidelines
 - Study “database normalization”
- But let’s stick to few basic ones
 - Each table should have a “primary key”
 - Must be unique on each row
 - Often id of type integer
 - Not part of actual data otherwise belonging to that data
 - Even if we have some kind of userid (perhaps email) for each person it should not be used as primary key

Designing databases, foreign key

- Foreign key is used to describe one-to-many relation
 - Or many-to-one, which ever way you want to look at it
 - Company may have several employees, group of employees belongs to the same company
- The “many end” of relation holds the foreign key
 - An extra field that has a value of primary key that identifies the item at “one end” of relation
 - Employee table would hold `company_id` that refers to id of the company

Designing databases, many-to-many

- Purchase has several products, single product may be used on several purchases
- A separate “binding table” is needed

| Product | |
|---------|--------|
| id | Name |
| 1 | Cheese |
| 2 | Milk |

| Purchase | |
|----------|----------|
| id | Customer |
| 1 | John |
| 2 | Mike |

Cheese was
bought by John
and Mike

| Id | Product_id | Purchase_id |
|----|------------|-------------|
| 1 | 1 | 1 |
| 2 | 1 | 2 |
| 3 | 2 | 2 |

Mike bought
cheese and milk

Some other rules

- Columns should be undividable
 - Name is in most cases a bad idea, use firstname, lastname instead
- Columns should not repeat
 - No extra1, extra2, extra3, use extrainfo-table with foreign key to original table instead
- Naming
 - Choose a naming convention and stick to it
 - Avoid camel casing, not all RDBMS:s support it
 - Use all lower case
 - Table names should be in singular form
 - company, person NOT companies, persons

Datatypes

- Each column in the table has a name and datatype
 - Like variables at program code
- The types available vary slightly between database products. Commonly available are
 - int, float and double
 - varchar(32), where 32 is the length of the field, the actual length may be up to 32 characters
 - char(6), where 6 is again the length, the string always has 6 characters
 - date and datetime
- For mysql you may study
 - <https://dev.mysql.com/doc/refman/8.0/en/data-types.html>

Null

- The value in a column may be null
 - When creating the table you can describe if the null value is allowed for specific column
- The behaviour of null varies slightly between RDBM:s
- Null testing is most often done with “is” instead of “=” in sql
- There are also database specific function that can be used in queries to replace a null value with empty string for example
 - `select ifnull(email,'not available') from person;`

SQL Language

Most important SQL statements

- Select – Query data
 - We'll mostly work with queries
- DML, these are rather straightforward, we'll just take a few examples
 - Insert
 - Update
 - Delete
- DDL
 - Create (database, table....)
 - Alter (database, table....)
 - Drop (database, table....)

SELECT basics

- Select basic syntax

```
SELECT      [column list]
FROM        [table list]
WHERE       [criteria]
```

```
GROUP BY   [column list]
ORDER BY   [column list];
```

Optional

- Get all data from customer-table

```
SELECT * FROM customer;
```

- Get name and price of all products

```
SELECT name,price FROM product;
```

- Get name and price for products where there are over 1000 items in stock

```
SELECT name,price
FROM product
WHERE in_stock > 1000
ORDER BY price desc,name;
```

???

Few notes

- Asterisk (*) is a wild card in column list = selects all columns
- You may define alias names for columns
 - Especially important if functions or calculations are used in column list
- Distinct makes sure all the rows returned are unique
- Tables may also have alias names
 - Important later on as we query data from several tables
- SQL keywords are not case sensitive but depending on the database engine the table – and column names may be

```
SELECT distinct p.id,  
       concat(p.firstname,' ',p.lastname) as name  
       salary*1.1 as increased_salary  
from person p;
```

Where criteria

The search criteria may become quite complex. It may hold

- Logical operators
`=, <>, !=`
- Boolean operators
`<, >, <=, >=`
- NULL –testing
`[NOT] LIKE`
- BETWEEN...AND –operators
`[NOT] EXISTS`
- Complex text-searches
`[NOT] IN`
- IN –operator
`[NOT] BETWEEN ... AND`
`IS [NOT] NULL`

Equality

- All females

```
SELECT * from person  
WHERE gender='F';
```

- All with specific last name

```
SELECT * from person  
WHERE lastname='Jones';
```

- Specific lastname but no woment

```
SELECT * from person  
WHERE lastname='Jones' AND gender<>'F';
```

- Lastname begins with something, firstname begins with T followed by any character and m

```
SELECT * from person  
WHERE lastname like 'Jo%' AND firstname like 'T_m';
```

IS NULL -operator

- Null value operates differently on different DBMS:s
 - Sorting: either first or last
 - Testing equality: always matches or never matches
- But is null / is not null works consistently between DBMS:s

```
SELECT    *  from person
WHERE     salary IS NULL;
```

```
SELECT    *  from person
WHERE     salary IS NOT NULL;
```

BETWEEN...AND –operator

- In most cases you could use $a \geq b$ AND $c \leq d$ also
- People whose salary is between 5000 and 5500

```
SELECT * FROM person  
WHERE salary BETWEEN 5000 AND 5500;
```

- Less than 5000 or greater than 5500

```
SELECT * FROM person  
WHERE salary NOT BETWEEN 5000 AND 5500;
```

IN –operator

- Is the value in a list
 - List is often a subquery

```
SELECT * FROM sales
WHERE salesdate IN (
    '01.12.2011', '01.01.2012',
    '01.02.2012', '01.03.2012'
);
```

- All persons that haven't bought something

```
SELECT * FROM person
WHERE id NOT IN (select personid from sales);
```

Joining data from several tables

- If table list has several tables by default all possible combinations are produced
 - WHERE clause must be used to limit the result
- Most often foreign keys are used to find matching values from other tables
 - But other information may be used also
- To consider
 - Do we want only to get information if the join criteria is satisfied
 - Inner join
 - Do we want to get all the information from one of the tables even if there is no matching data on second table
 - Outer join

Inner join

- Most often used join type
- Only returns rows where there is a binding
- Get all data from person table + the name of the company but only for people for whom the company is specified

```
SELECT p.*, c.name  
FROM person p, company c  
WHERE p.companyid=c.id
```

left (outer) join

- Get all information from the table specified in the FROM clause
 - Even if the data doesn't exist on the joined table
- You can limit the final result with WHERE-clause after the joins
- Everything from person table, c.name is null if company is not specified for the person

```
SELECT p.*, c.name FROM person p  
LEFT JOIN company c ON (c.id=p.companyid)
```

Right (outer) join

- Instead of left join you can also use right join
- Works the other way around
 - Everything from the joined table even if there is no match on the table specified in the FROM-table
- Recommendation is to use left joins when possible
 - They are easier to understand

```
SELECT p.* , c.name FROM company c  
RIGHT JOIN person p ON (p.companyid=c.id)
```

Some functions

- Functions are always DBMS-specific
 - All have similar functions but their names and parameter lists may differ
 - Study the documentation of DBMS you end up working with
 - For MySQL
<https://dev.mysql.com/doc/refman/8.0/en/functions.html>
- Few examples
 - String concatenation: `concat(firstname, ' ', lastname)`
 - Replace null value: `ifnull(email, 'not specified')`
 - Date part of current timestamp: `date(now())`

Group by

- Most commonly group by –clause is used in association with aggregate function
 - Sum, Avg, Count
- You must specify all columns not used in aggregate function to group by list of columns

- How many employees company has

```
select c.name, count(p.id) as empl_count
from company c, person p
where p.company_id=c.id
group by c.name;
```

Subqueries and Views

Subquery

- You use a separate query as a part of original query
 - You can use the result of subquery as table (alias name must be specified) in the FROM or JOIN list of tables
 - You can use the result of single column query with WHERE ... IN (...) -clause

Products are produced on different locations, our departments also work on different locations. Couple of ways to use this correlation

```
SELECT d.name as dname, p.name as pname
FROM product p, (SELECT name, location FROM department) d
WHERE d.location=p.location;
```

```
SELECT * FROM product
WHERE location in (SELECT location FROM department)
```

Views

- Views are predefined queries that may be used as tables on other queries
 - They don't store data
 - They just help build queries that would otherwise contain several subqueries
- View is created with `CREATE VIEW name_of_view AS SELECT` Command
 - After that you can use the view in other queries as it was a table
 - Some DBMS:s allow you to modify data through view as long as it uniquely identifies the underlying table

Data validation, constraints

- Table declarations may include validation rules. Invalid updates will not go through.
- Of those rules we just use
 - Primary key
 - Foreign key
 - not null for column (null value is not allowed)
- But there are more ways to validate data
 - Check <https://dev.mysql.com/doc/refman/8.0/en/create-table-check-constraints.html>

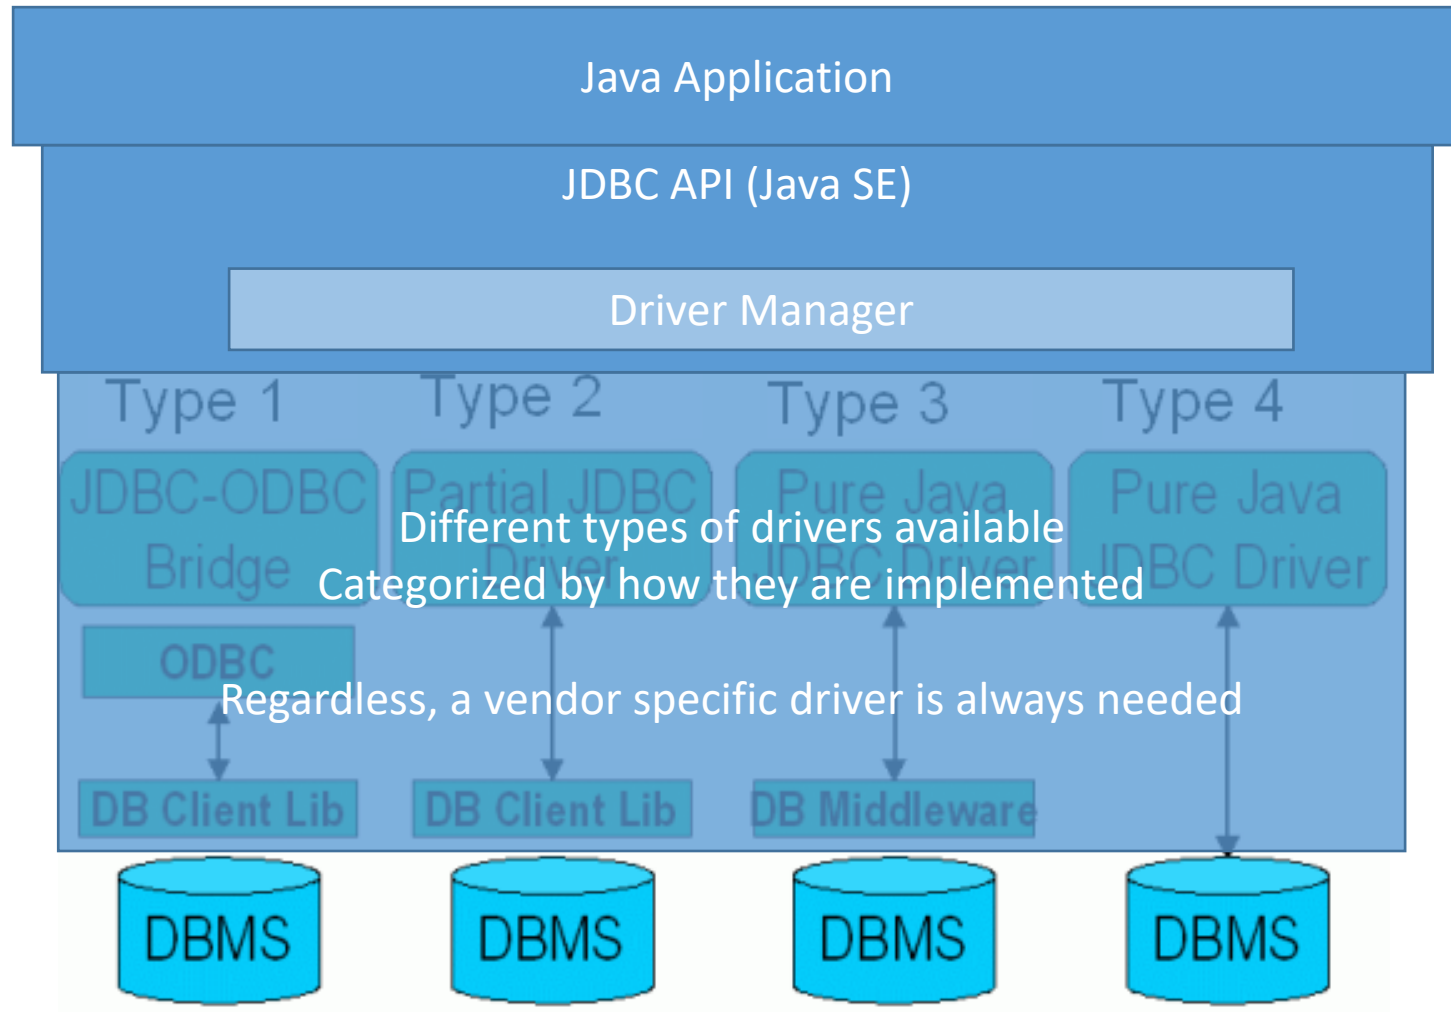
JDBC

Connecting to databases from Java application

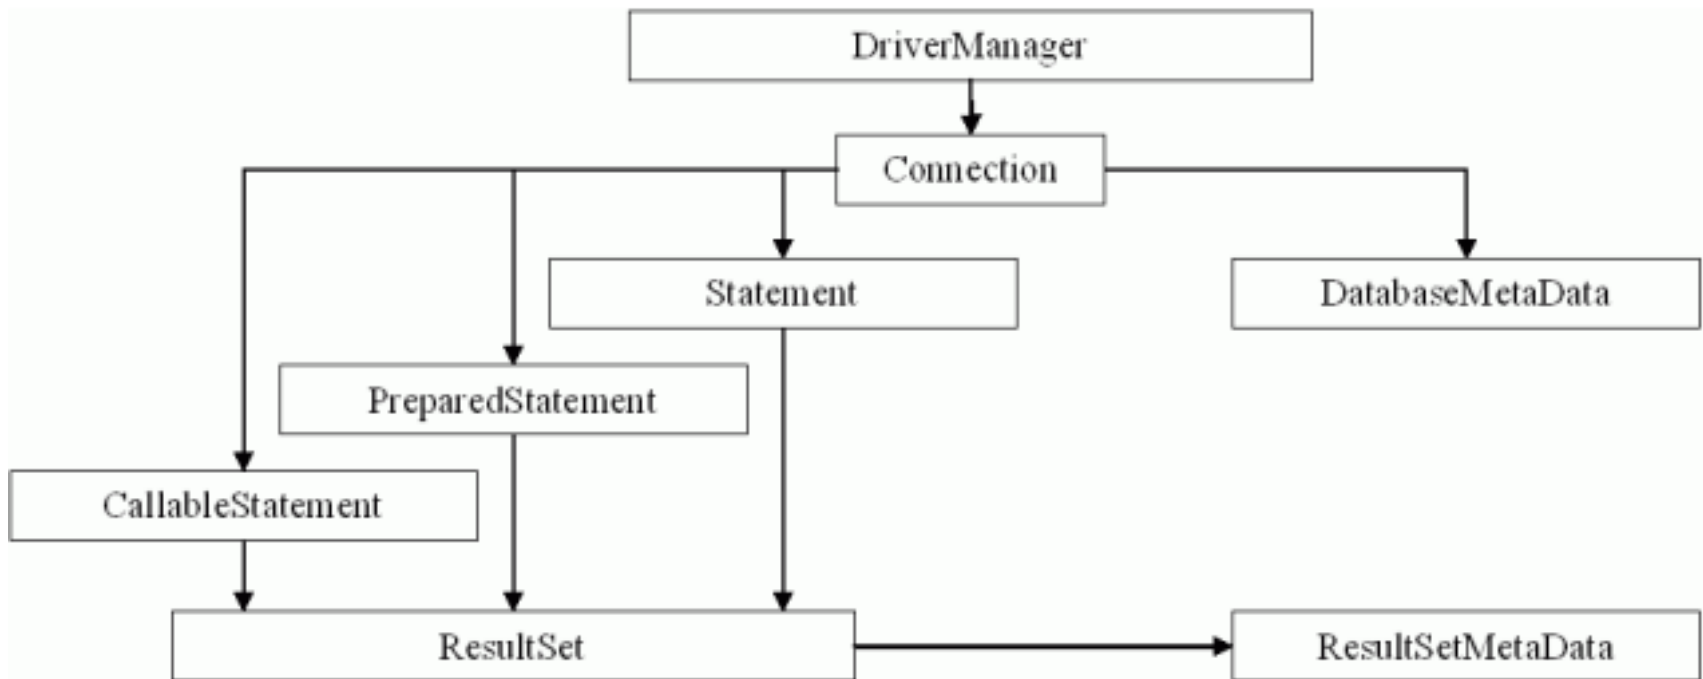
JDBC

- Java Database Connectivity → JDBC
 - A standard for database-independent connectivity between the Java programming language and a wide range of databases
 - An interface for accessing database management systems (DBMS)
- "Write Once, Run Anywhere" capabilities for applications that require access to enterprise data
 - No vendor-specific database (SQL) queries
 - Based on X/Open SQL CLI (like ODBC).
- Wide range of data sources possible:
 - SQL databases
 - Other tabular data sources, such as spreadsheets or flat files
- The JDBC API makes it possible to do three things:
 - Establish a connection with a database / tabular data source
 - Execute SQL statements
 - Process the results

JDBC Architecture



JDBC Object Model



Establishing a Connection

- Two steps: loading the driver and making the connection

1. Making sure the driver is available

- Not an absolutely necessary step, but may help to pinpoint why the initial connection fails
- Easy, just one line of code that basically loads the driver into memory:

```
Class.forName("com.mysql.jdbc.Driver");
```

- Driver class is different for each database
 - Driver is usually implemented by the database provider
 - Check from your database documentation which class name to use

Establishing a Connection

2. Making the connection

- The JDBC `DriverManager` class manages the JDBC drivers that are installed on the system
- It also provides factory functions for creating connections using the chosen driver

```
Connection conn = DriverManager.getConnection(  
    "db url", "loginName", password);
```

- Note that the database url depends on the driver used
- For MySQL:
`jdbc:mysql://localhost:3306/myJdbcDB`

Establishing a Connection

```
Connection conn = null;
try {
    Class.forName("com.mysql.jdbc.Driver");
    String url = "jdbc:mysql://localhost:3306/myJdbcTest";
    conn = DriverManager.getConnection(url , "myUser", "myPw");

    // ...use connection...

    // Remember to close the connection when done using it!

    conn.close();          // Might throw SQLException - omitted here

} catch (ClassNotFoundException cne) {
    System.out.println("Couldn't find driver!");
} catch (SQLException se) {
    System.out.println("Couldn't establish connection");
}
```


Exercise

- Copy instructor pointed DBCreator.java to a location of your choosing
 - Note the package name in the file
- Compile the code from command line
- Run the code
- You need the driver, copy the instructor pointed jar to your “project” folder
- Run the code
 - `java -cp .;\xxxx.jar jdbc.DBCreator`

Creating JDBC Statements

- A `Statement` object is what sends your SQL statement to the DBMS.
 - Simply create a `Statement` object using your `Connection` object and then execute it

```
Statement stmt = conn.createStatement();  
String sql = "INSERT INTO table1 (one, two) VALUES ('one', 'two')"  
stmt.executeUpdate(sql);
```

- For SQL statements that create or modify tables, the method to use is `executeUpdate()`
- For a `SELECT` statement, the method to execute the query is `executeQuery()`

Creating JDBC Prepared Statements

- The `PreparedStatement` object contains a pre-compiled SQL statement
 - When executed, the DBMS can run the `PreparedStatement`'s SQL statement without having to compile it first
 - Much faster for statements that get executed multiple times than using regular `Statement` objects
 - Easier to avoid SQL injection attacks as inserted values are automatically escaped

```
String sql = "UPDATE table1 set one = ?, two = ?";  
PreparedStatement prepStmt = conn.prepareStatement(sql);  
prepStmt.setInt(1, 123);           // First ?-mark  
prepStmt.setString(2, "myNewValue2"); // Second ?-mark  
prepStmt.executeUpdate();
```

Exercise

- Modify DBCreator slightly
 - Use Visual Studio Code for this exercise
- Create a new company
- Query a company by id

JDBC datatypes

- Data types in SQL and Java are not identical
 - JDBC defines a standard mapping from the JDBC database types to Java types
 - For example, a JDBC `INTEGER` is normally mapped to a Java `int`
 - This enables a simple interface for reading and writing JDBC values as simple Java types
- JDBC defines a set of generic SQL type identifiers in the class `java.sql.Types`
 - These types have been designed to represent the most commonly used SQL types, for example:

`VARCHAR/CHAR/LONGVARCHAR`

`DATE`

`INTEGER`

`TIME`

`java.lang.String`

`java.sql.Date`

`int`

`java.sql.Time`

JDBC ResultSet

- The `executeQuery()` method returns a `ResultSet` object
 - Both for `Statement` and `PreparedStatement`
- Provides access to the result of the query – sort of like a table
 - Table rows are retrieved in sequence
 - A `ResultSet` maintains a cursor pointing to its current row of data
- The `next()` method is used to iterate through the rows of the tabular results
 - Cursor is initially positioned just above the first row of a `ResultSet` object
 - The first call to `next()` moves the cursor to the first row and makes it the current row
 - Call `next()` again to move to the next row

JDBC ResultSet Example

```
String sql = "SELECT * FROM customer";
Statement stmt = conn.createStatement();

// Execute query to get the customers in a ResultSet
ResultSet rs = stmt.executeQuery(sql);
String id, name, address;
// Iterate over the ResultSet using next()
while(rs.next()){
    // Use the appropriate getter functions to access columns
    id = rs.getString("id");
    name = rs.getString("name");
    address = rs.getString("address");
    System.out.println(id + " " + name + " " + address);
}

stmt.close();    // Closes also the ResultSet object
```

Updates

- You can always pass SQL INSERT, UPDATE and DELETE statements to `stm.executeUpdate`
- But, especially for updates you can also update the data in the result set

```
public void update(Person p){
    String q="SELECT * FROM Person where ID=?";
    try{
        Connection con=getConnection();
        PreparedStatement stm=con.prepareStatement(q,
            ResultSet.TYPE_FORWARD_ONLY,
            ResultSet.CONCUR_UPDATABLE);
        stm.setInt(1, p.getId());
        ResultSet rs=stm.executeQuery();
        if (rs.next()){
            rs.updateInt("company_id", p.getCompanyId());
            rs.updateString("name", p.getName());
            rs.updateString("email", p.getEmail());
            rs.updateDate("contact_date", p.getContactDate());
            rs.updateRow();
        }
        rs.close();
        stm.close();
        con.close();
    }
    catch(Exception ex){
        ex.printStackTrace();
    }
}
```


JDBC transactions

- When a connection is created, it is in auto-commit mode
 - Each individual SQL statement is treated as a transaction and will be automatically committed right after it is executed

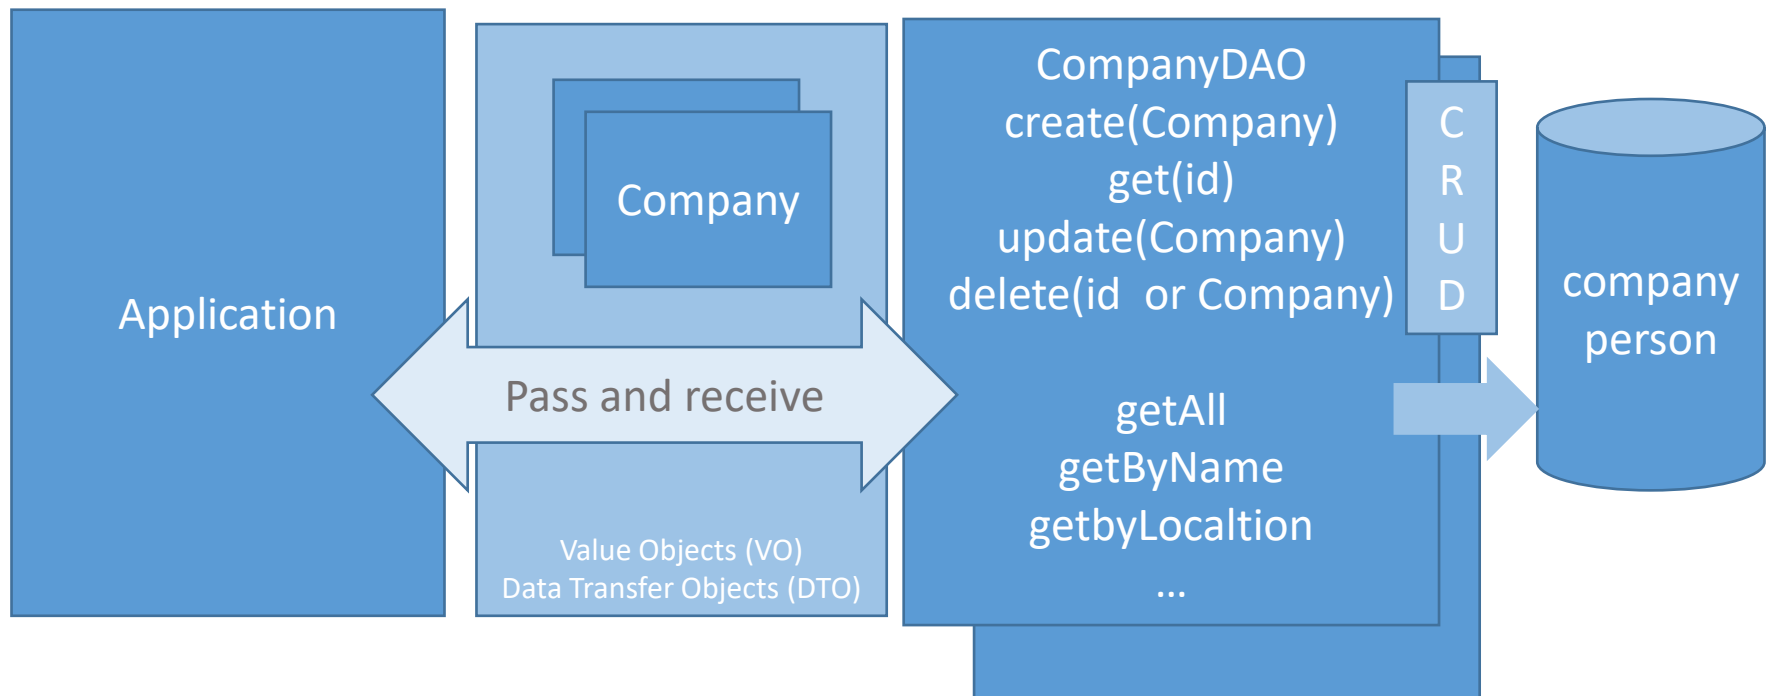
- Way to allow two or more statements to be grouped into a transaction is to disable auto-commit mode

```
Connection conn = ...;  
conn.setAutoCommit(false);
```

- Once auto-commit mode is disabled, no SQL statements will be committed until you call `commit()` explicitly
 - All statements executed after the previous call to `commit()` will be included in the current transaction and will be committed together as a unit

DAO-pattern

- Database has tables (company, person)
- We work with object of Person class and Company class
- Let's implement a class against each table that translates table data into objects



Exercise

- Create a new book application “BookJdbc”
- Add a runtime dependency to MySQL-connector
 - Right click dependencies on Project-view
 - Check the pom.xml after adding the dependency
- Start by creating Author-class
 - Same fields as you have at author table
- Then create AuthorDAO-class with
 - Connection getConnection();
 - Author authorFromResultset(ResultSet rs)
 - Author get(int id)
 - Uses the first two to return author that has been fetched from the database
- Test in the main

Exercise

- To the AuthorDAO implement
 - List<Author> getAll()
 - Author create(Author a)
 - The id must change into the id of newly created author
 - Author update(Author a)
 - Return author as it was saved to the database
 - boolean delete(int id)
- To the mainclass implement showAllAuthors
 - That calls AuthorDAO.getAll and then displays all authors
 - This method just helps your testing

Exercise

- Create BookDAO
 - It should hold the CRUD methods as AuthorDAO did but it is not necessary to implement all those for this exercise
 - If you have time, you may do that, otherwise just add a comment that they should be there
- getBooksOfAuthor(int authorid)
 - Does this belong to BookDAO or AuthorDAO?
- getAuthorOfBook(int book)
 - Does this belong to BookDAO or AuthorDAO
- getBooksWithAuthor()
 - Returns some information about book and at least the name of the author

Exercise

- To your Book-class add Author author –field
- To you Author-class add ArrayList<Book> books –field
- To the BookDAO add void getAuthor(Book book)
 - That reads the author of the book and places it to the author-field
- To the AuthorDAO add void getBooks(Author author)
 - That reads the books from the author and places them to books-field of author object

Exercise

- LoanDAO
 - What methods should go in there?